

Model Generalization

(Draft)

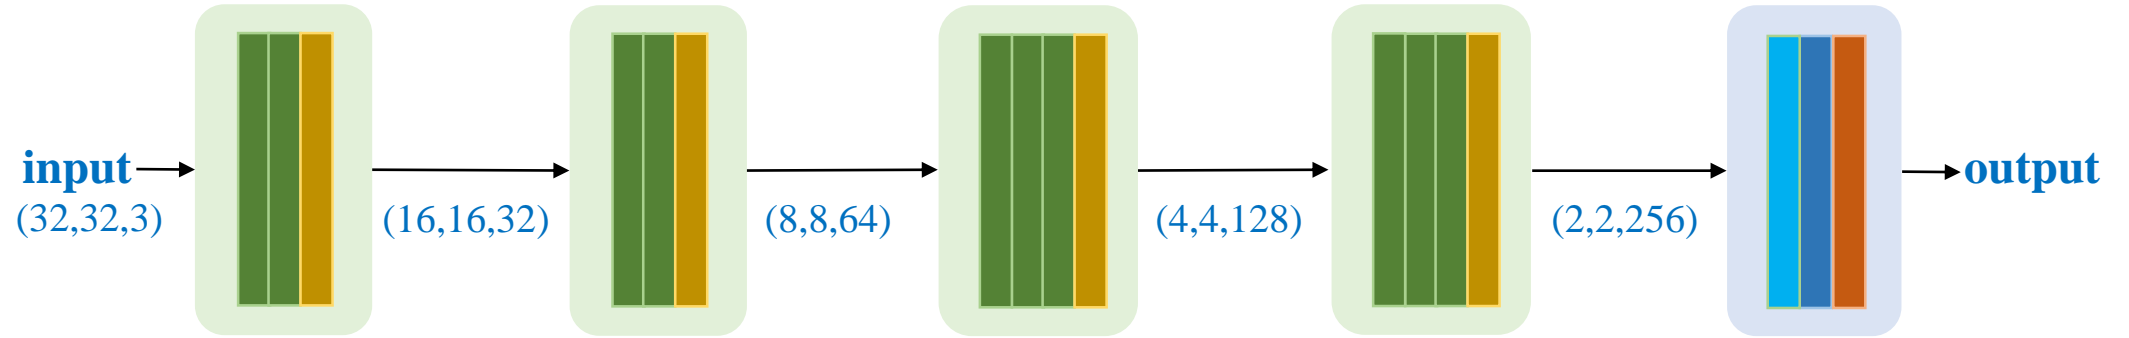
Quang-Vinh Dinh
Ph.D. in Computer Science

Outline

- **Introduction to Numpy**
- **Numpy Array Indexing**
- **Numpy Array Operations**
- **Broadcasting**
- **Data Processing**

Model Generalization

Cifar-10
Dataset



Data Normalization
(convert to 0-mean
and 1-deviation)

$$\bar{X} = \frac{X - \mu}{\sigma}$$

$$\mu = \frac{1}{n} \sum_i X_i$$

$$\sigma = \sqrt{\frac{1}{n} \sum_i (X_i - \mu)^2}$$

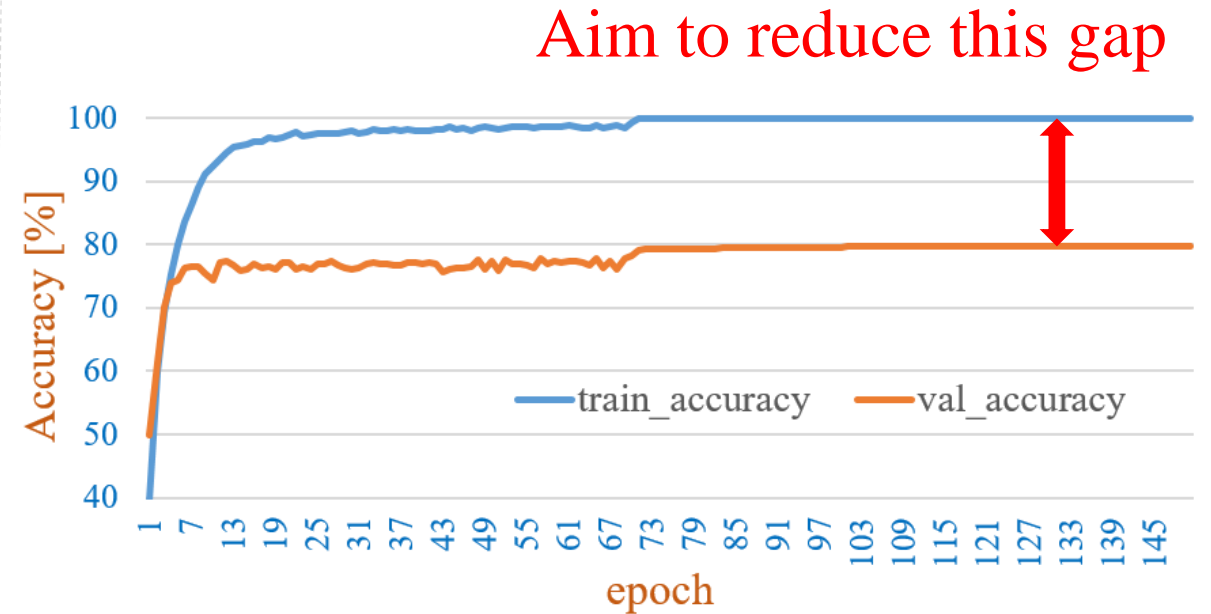
(3x3) Convolution
padding='same'
stride=1 + ReLU

(2x2) max pooling

Flatten

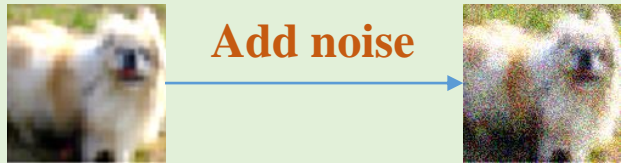
Dense Layer-10
+ Softmax

Dense Layer-512
+ ReLU



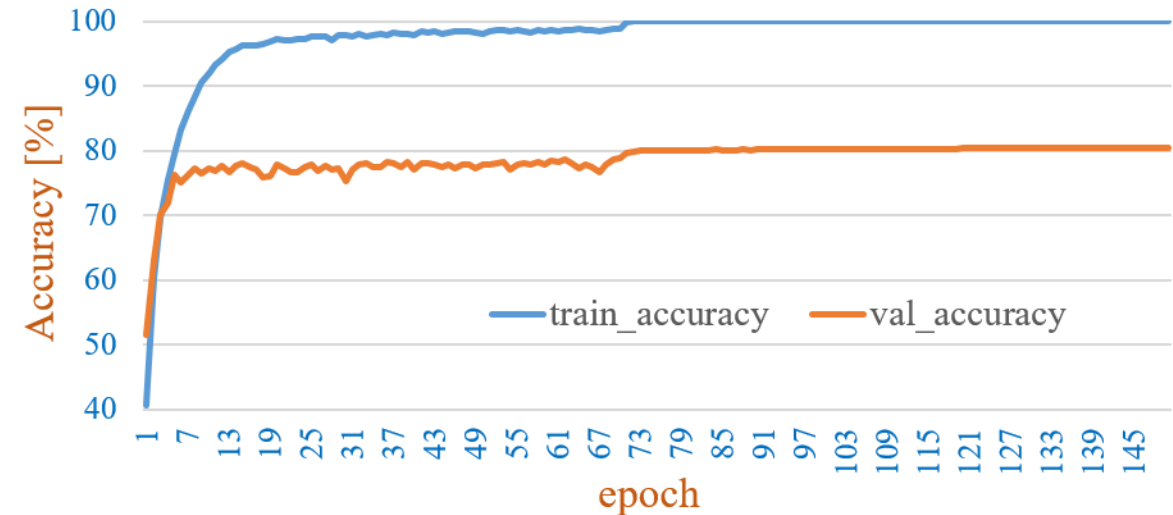
Model Generalization

❖ **Trick 1: 'Learn hard, ' – randomly add noise to training data**



In Keras

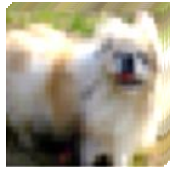
```
1 if tf.random.uniform(()) > 0.5:
2     noise = tf.random.normal((32, 32, 3))/100.0
3     image = image+noise
4
5     return image, label
```



**val_accuracy increases
from ~80.2% to ~80.9%**

Model Generalization

❖ Trick 2: Batch normalization



mini-batch 1



mini-batch 2

$$(\mu_1, \sigma_1) \neq (\mu_2, \sigma_2)$$

very
likely



Add noise to the output of BN layers

Input data for a node in batch normalization layer

$$X = \{X_1, \dots, X_m\}$$

m is mini-batch size

Compute mean and variance

$$\mu = \frac{1}{m} \sum_{i=1}^m X_i \quad \sigma^2 = \frac{1}{m} \sum_{i=1}^m (X_i - \mu)^2$$

Normalize X_i

$$\hat{X}_i = \frac{X_i - \mu}{\sqrt{\sigma^2 + \epsilon}}$$

ϵ is a very small value

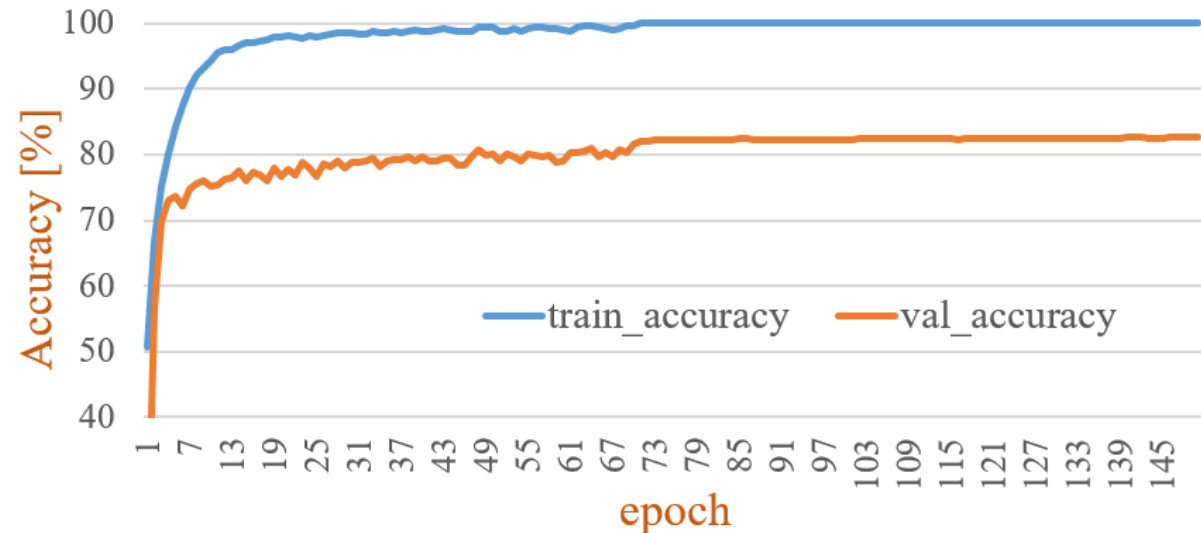
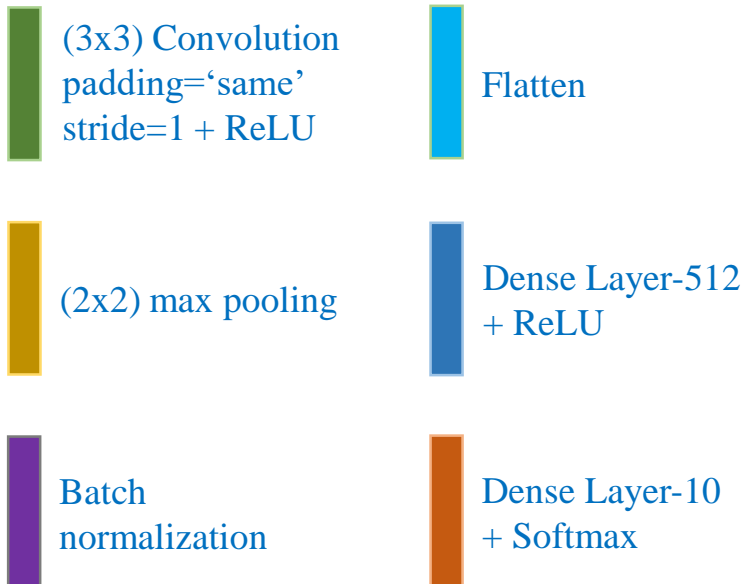
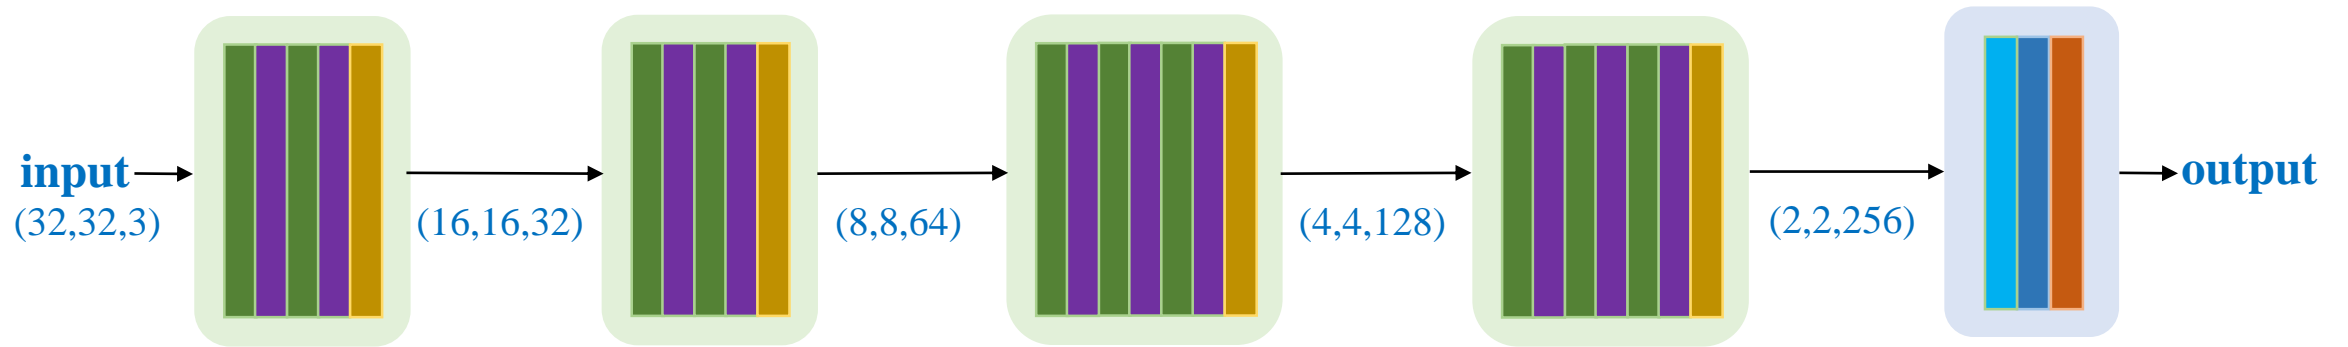
Scale and shift \hat{X}_i

$$Y_i = \gamma \hat{X}_i + \beta$$

γ and β are two learning parameters

Model Generalization

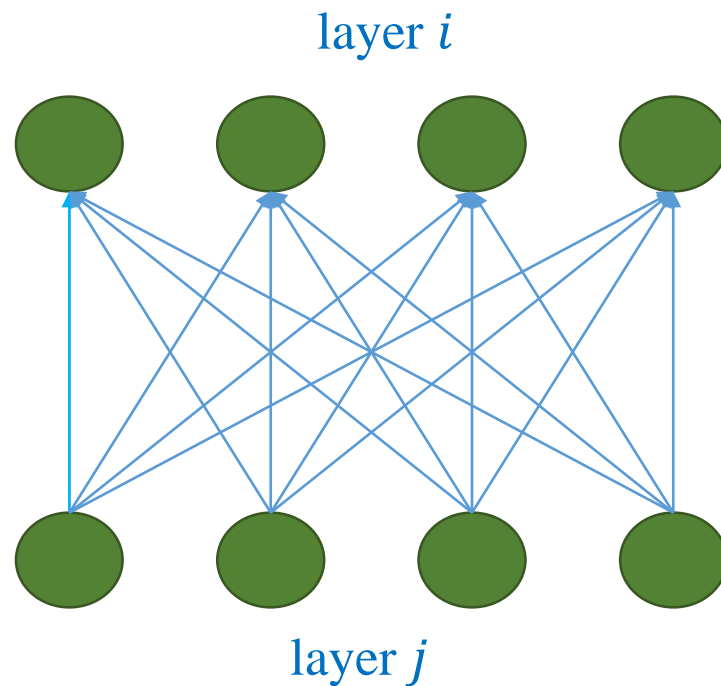
❖ Trick 2: Batch normalization



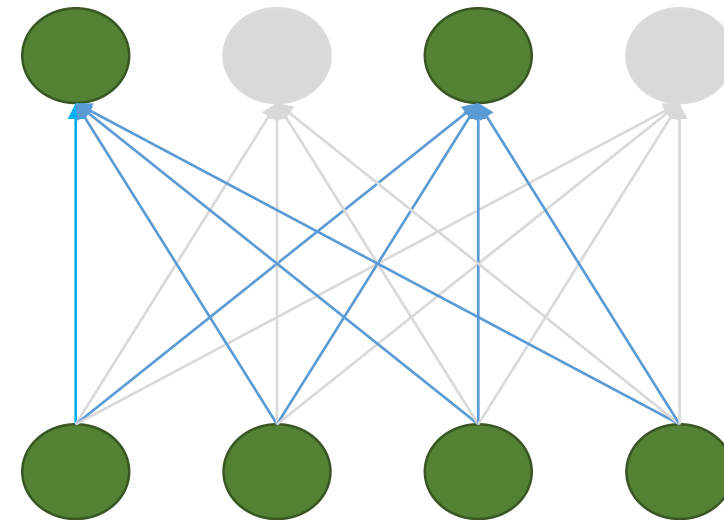
val_accuracy increases from ~80.9% to ~82%

Model Generalization

❖ Trick 3: Dropout



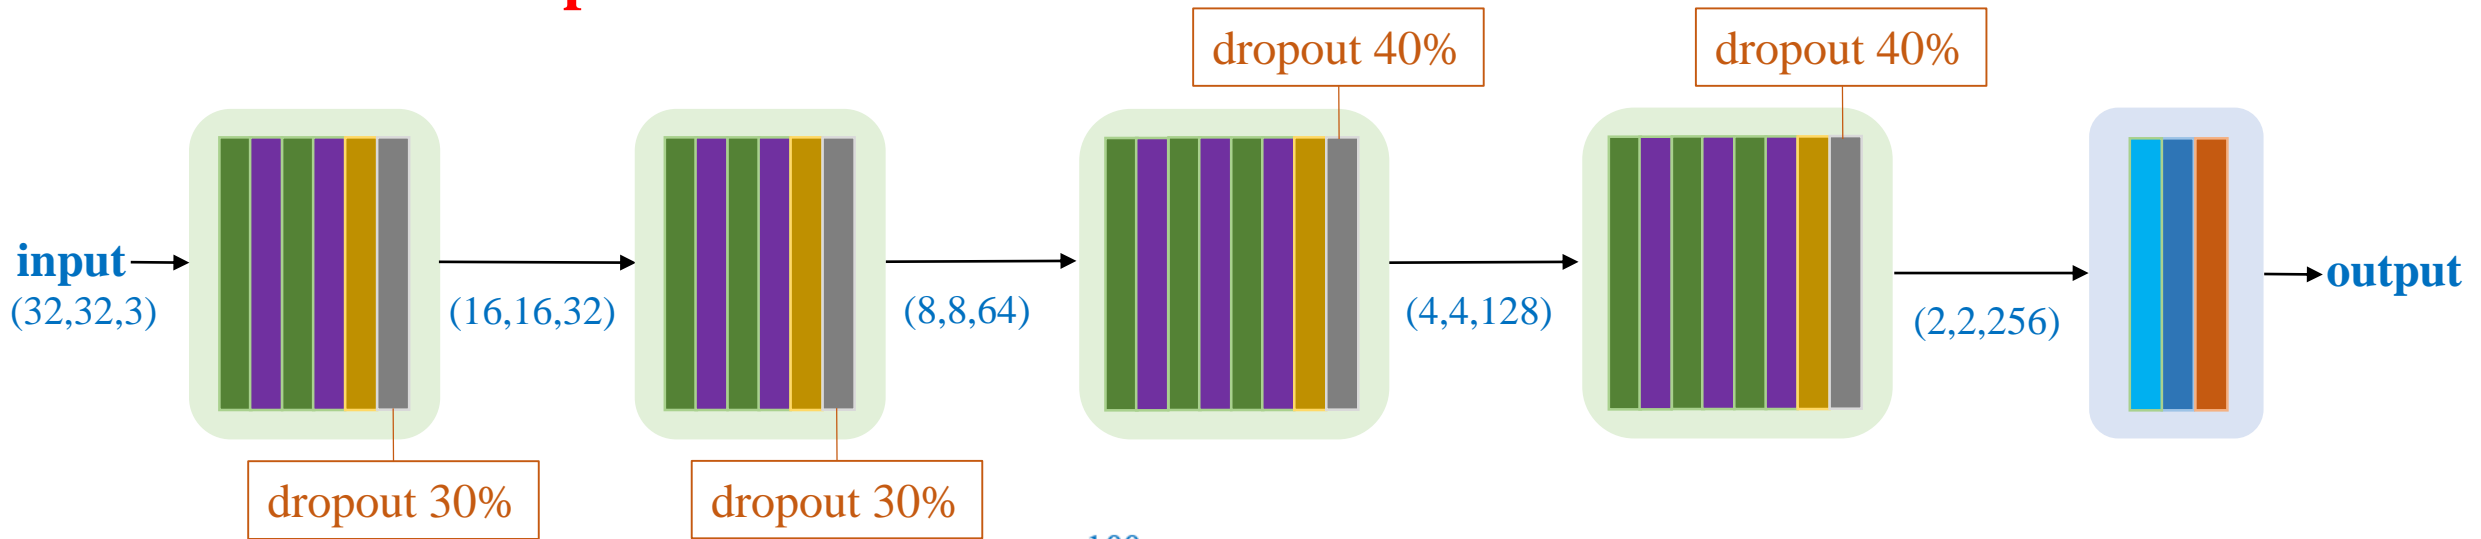
Apply dropout 50% to layer i



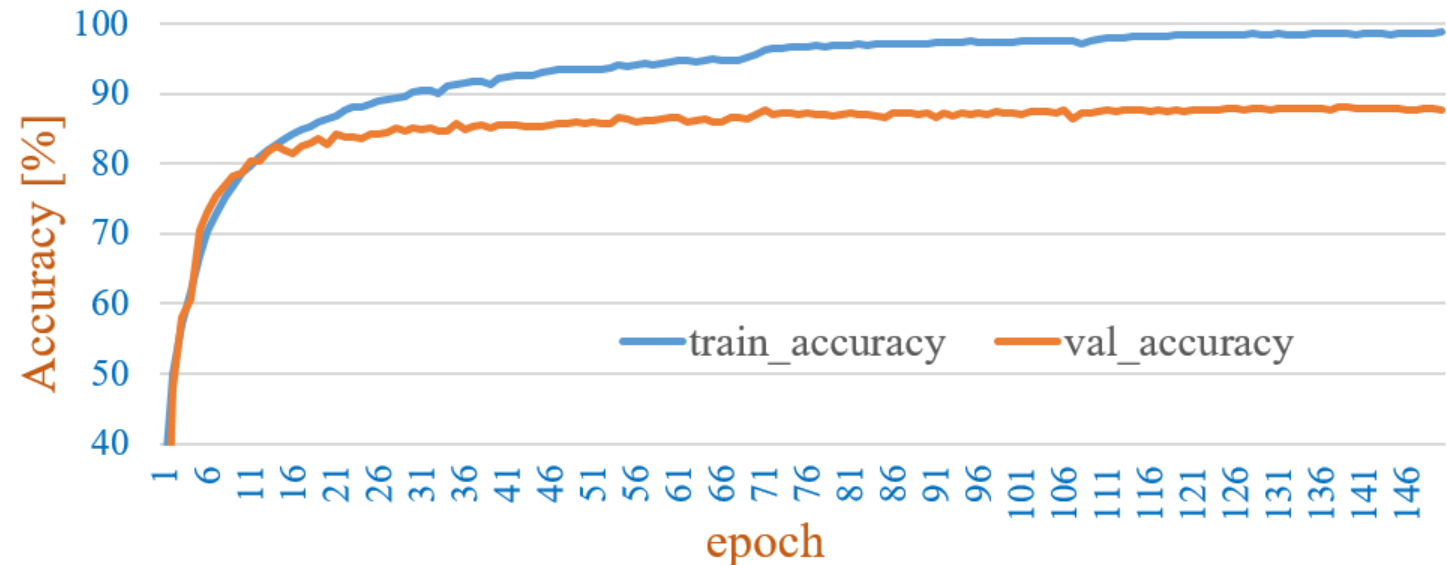
~50% nodes randomly selected in the i^{th} layer are set to zeros (kind of noise adding)

Model Generalization

❖ Trick 3: Dropout



**val_accuracy
increases from
~82% to ~87.9%**



Model Generalization

❖ Trick 4: Kernel regularization

$$L = \text{crossentropy} + \underbrace{\lambda_1 \|W\|}_{L_1 \text{ regularization}} + \underbrace{\lambda_2 \|W\|^2}_{L_2 \text{ regularization}}$$

Prevent network from
focusing on specific features

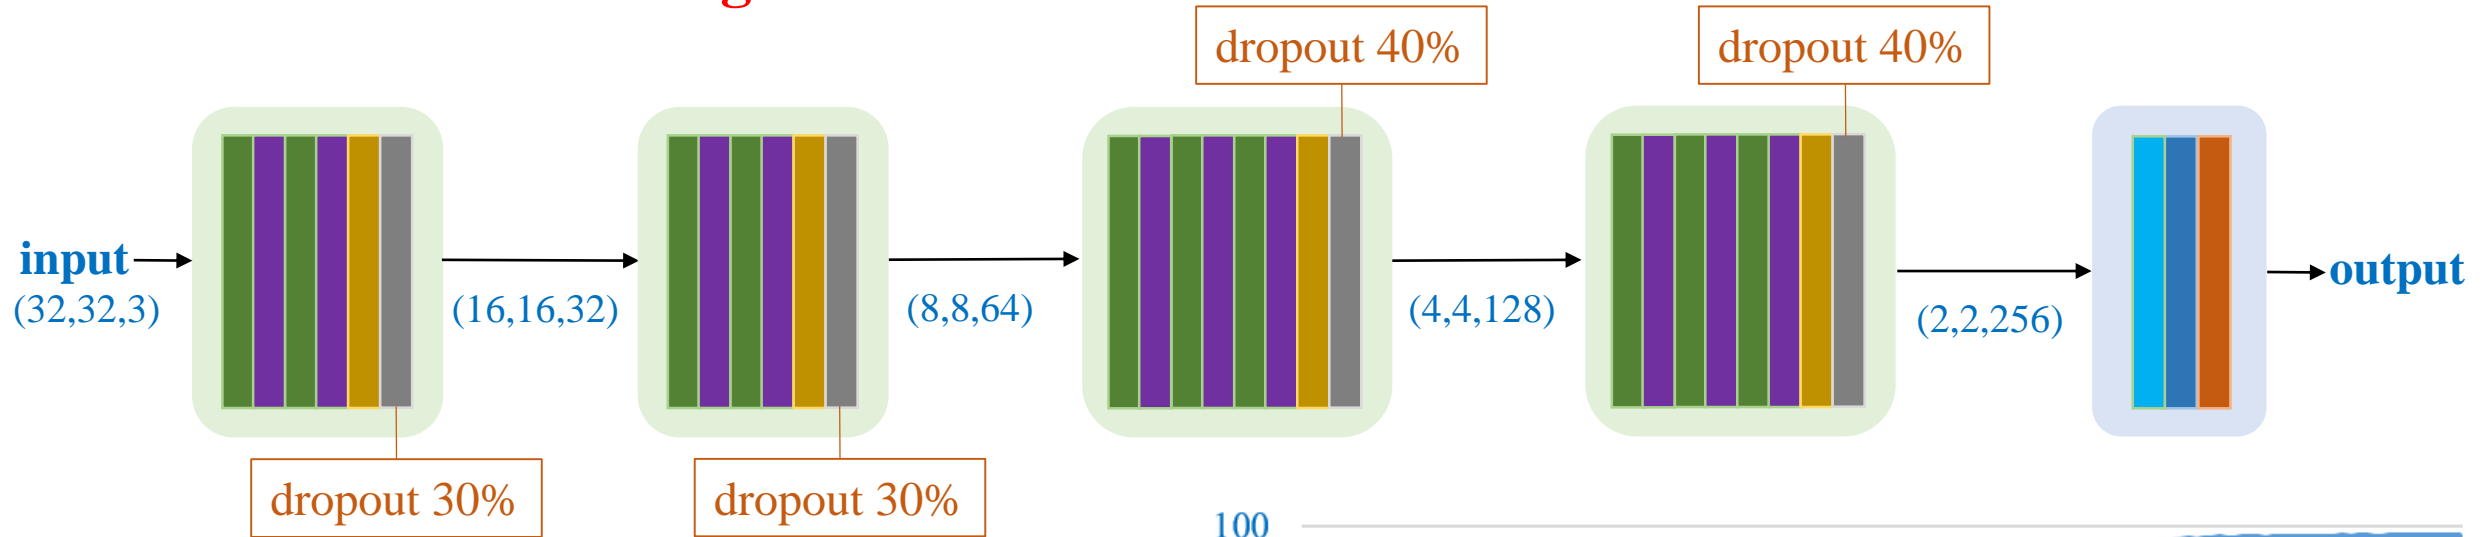
Smaller weights
→ simpler models

In keras

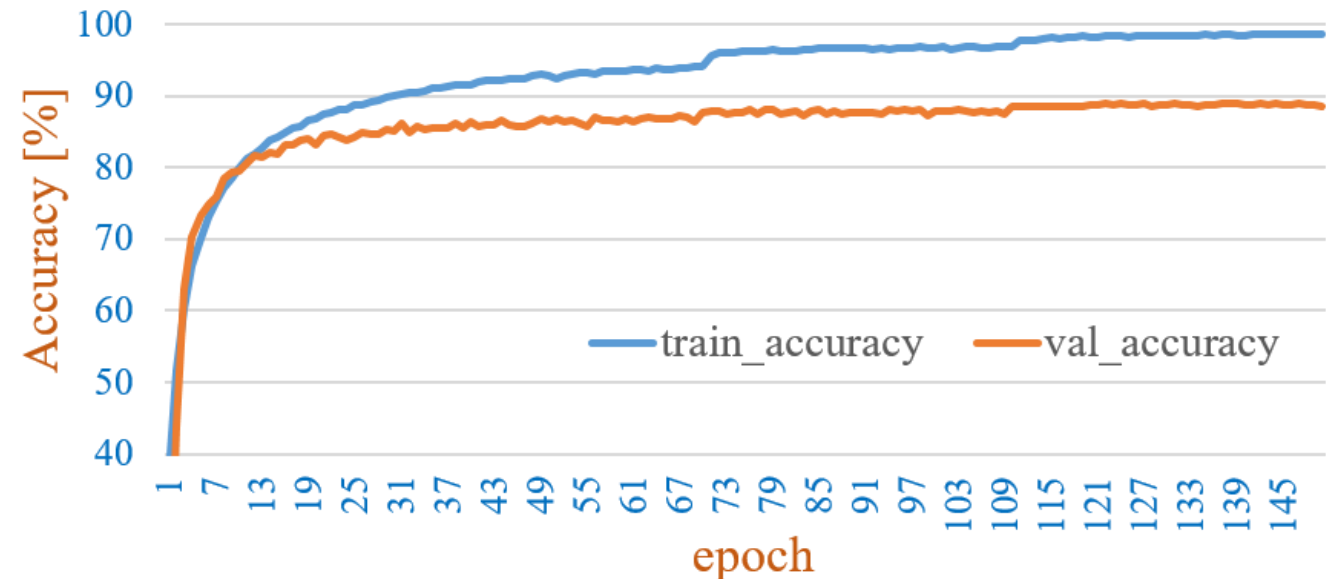
```
Conv2D(32, (3,3), padding='same', activation='relu',  
       kernel_regularizer=regularizers.l1_l2(decay1,decay2))
```

Model Generalization

❖ Trick 4: Kernel regularizer



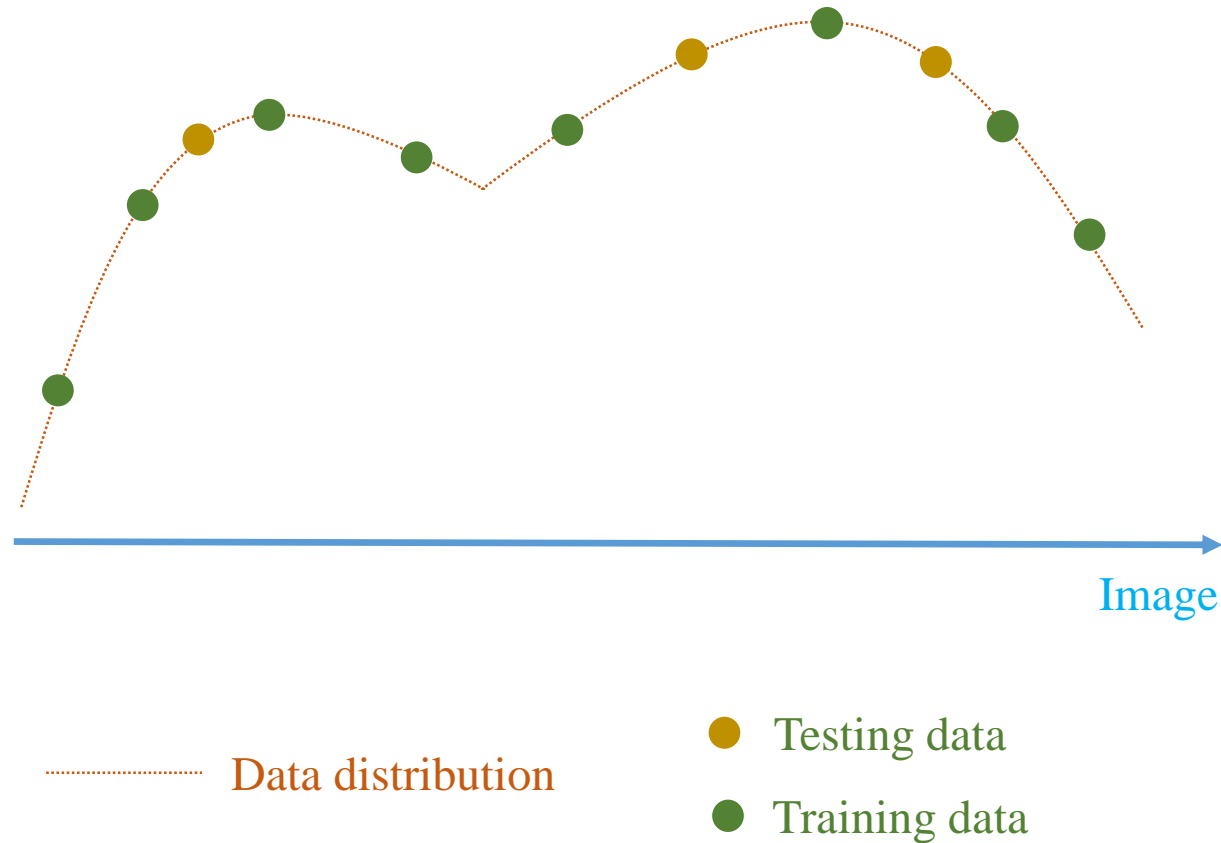
**val_accuracy
increases from
~87.9% to ~88.9%**



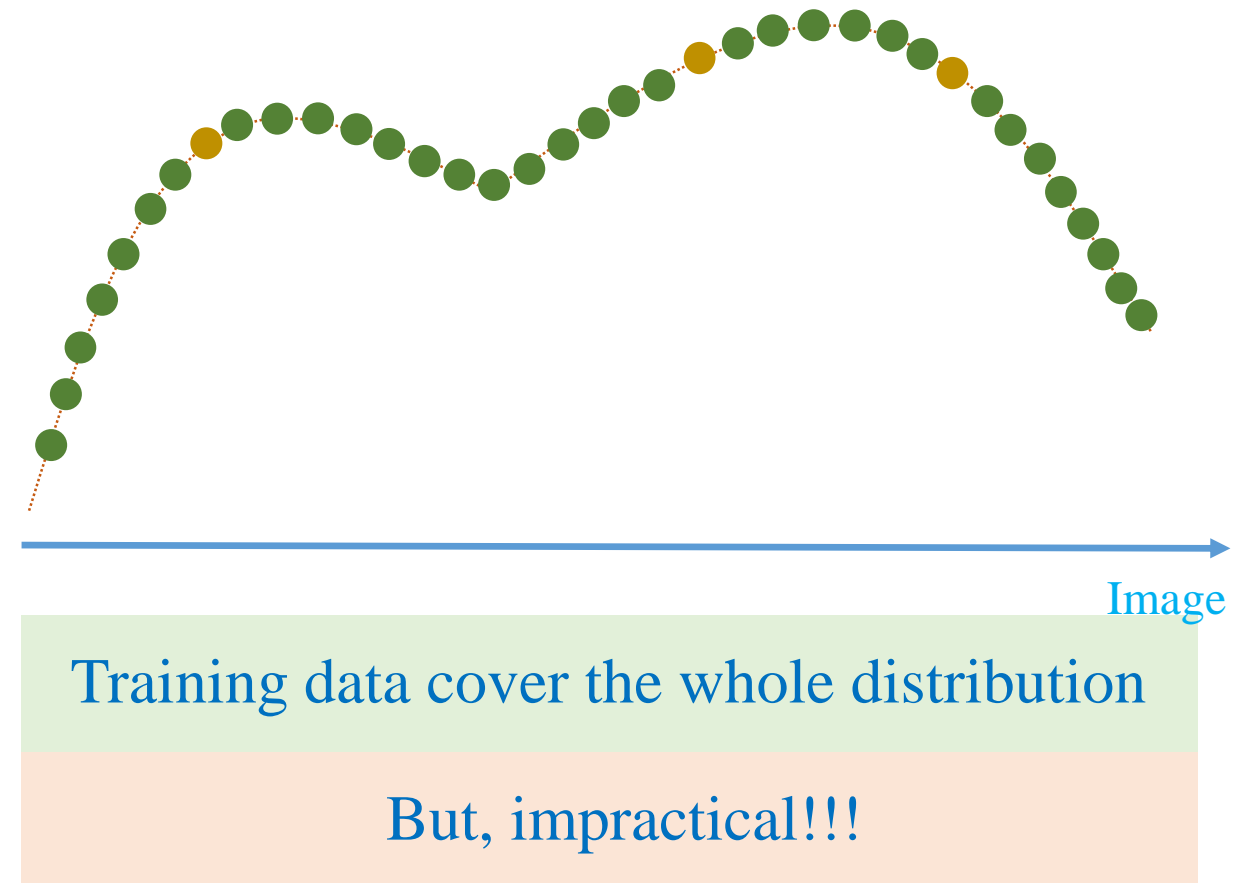
Model Generalization

❖ Trick 5: Data augmentation

A normal case

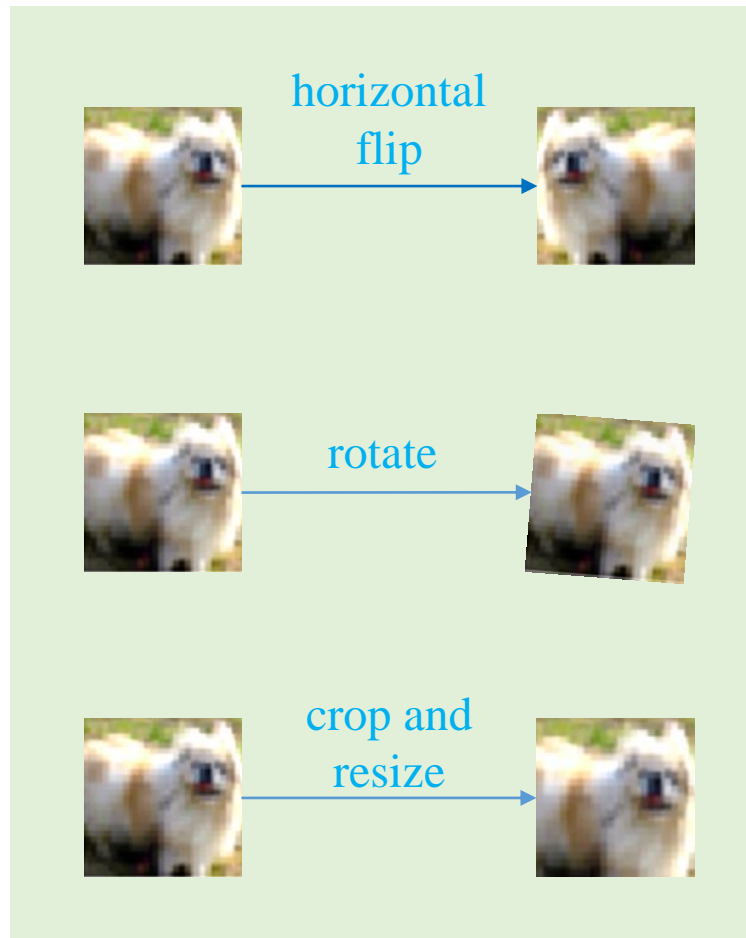


A perfect case: Have unlimited training

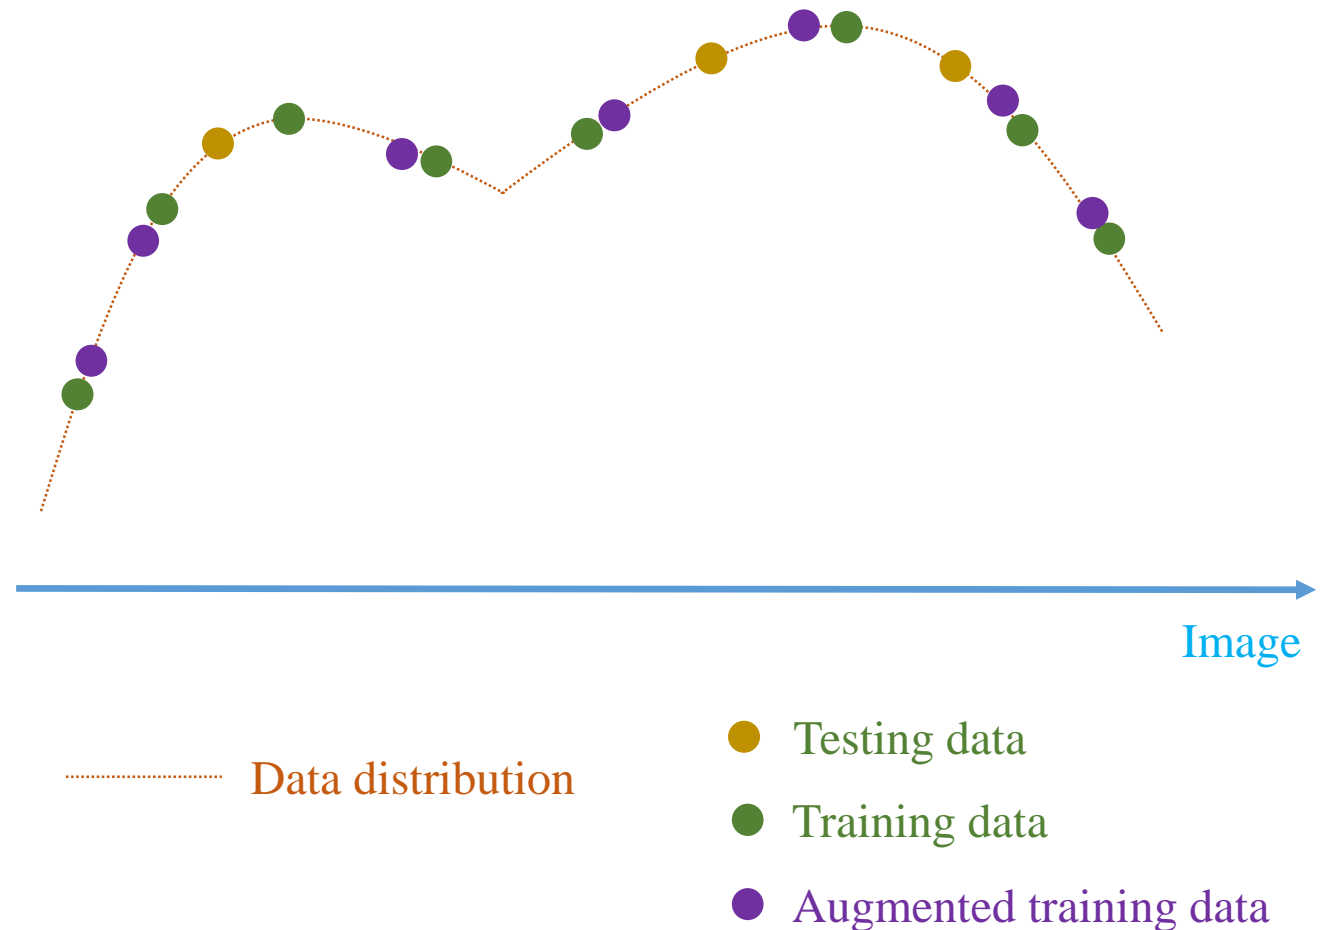


Model Generalization

❖ Trick 5: Data augmentation



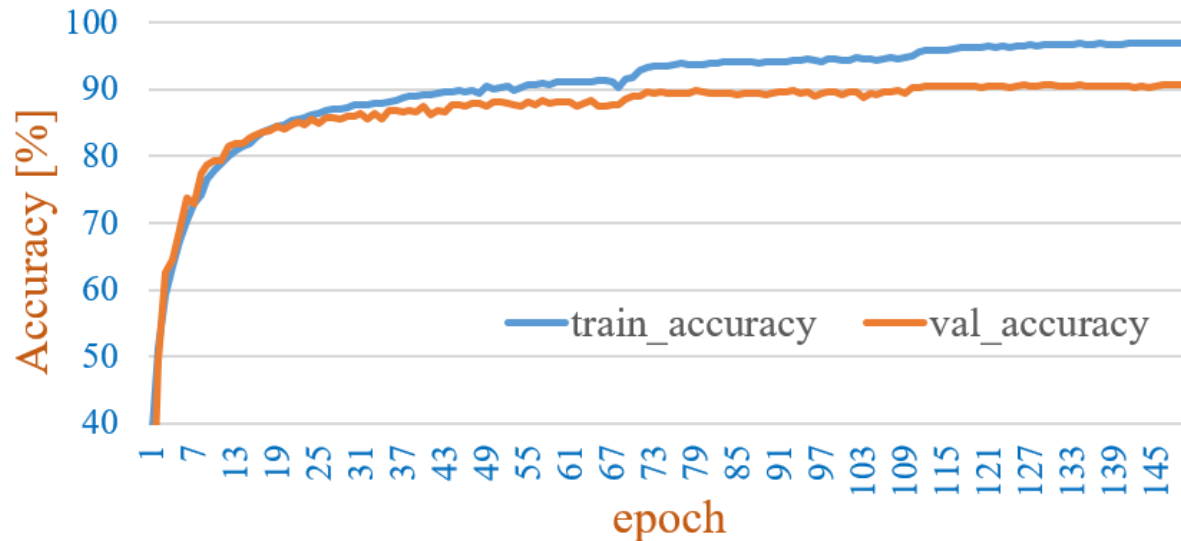
Increase data by altering the training data



Model Generalization

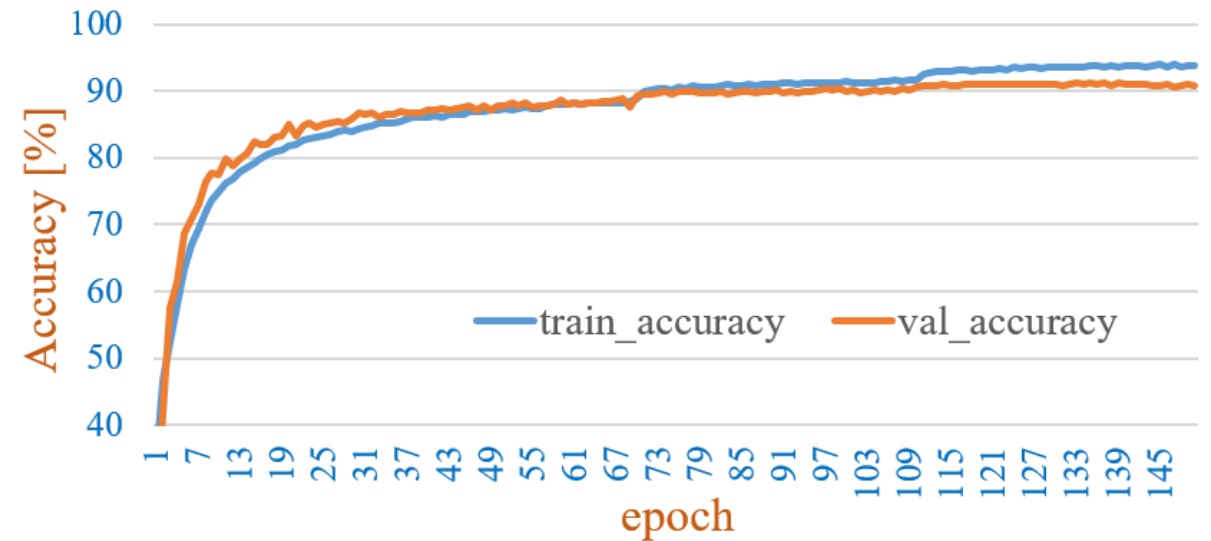
❖ Trick 5: Data augmentation

Horizontal flip



val_accuracy reaches to ~90.7%

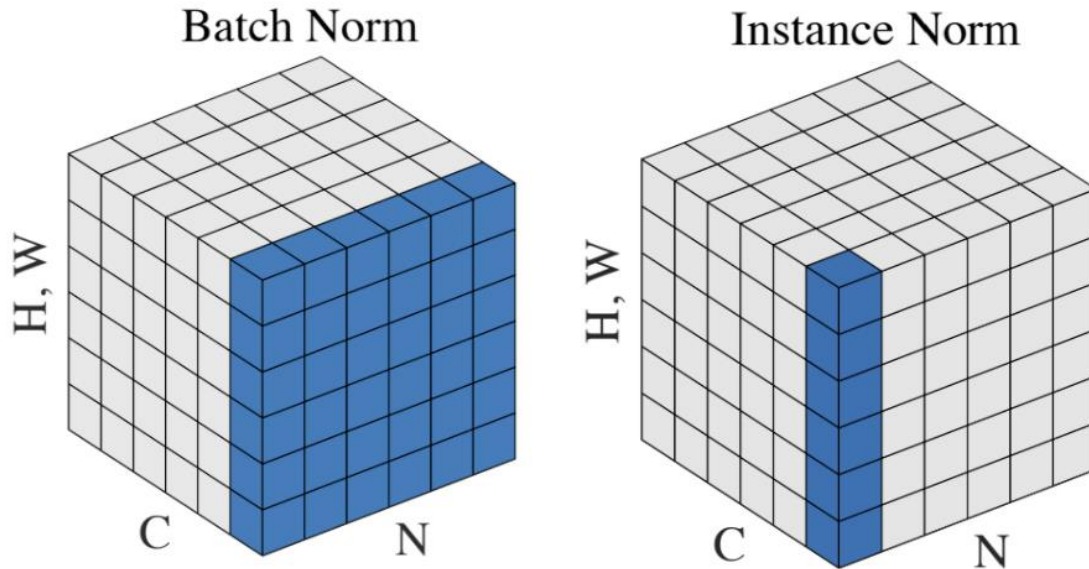
Horizontal flip + crop-and-resize



val_accuracy reaches to ~91.2%

Model Generalization

❖ Trick 6: Instance normalization



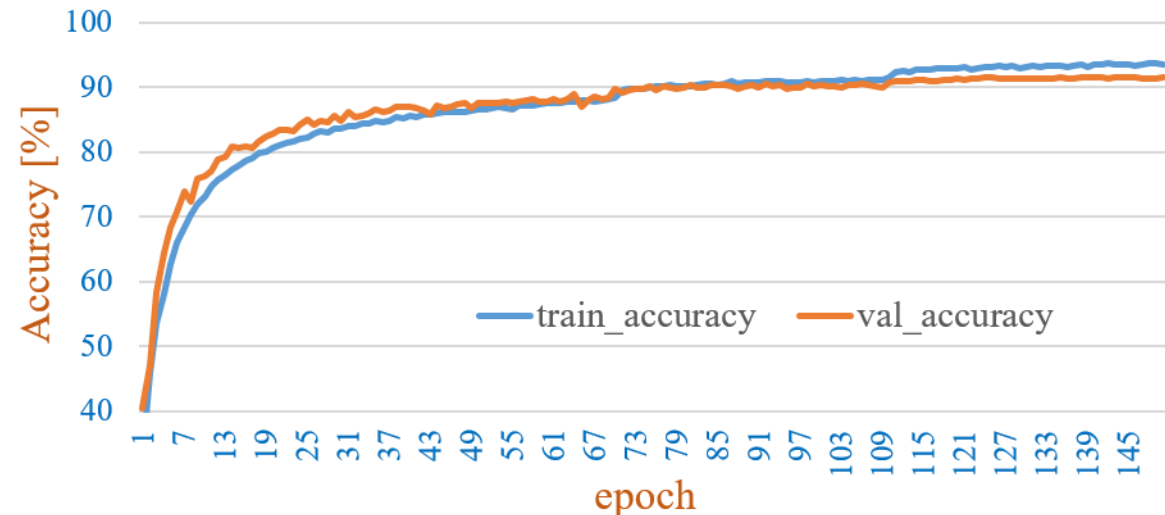
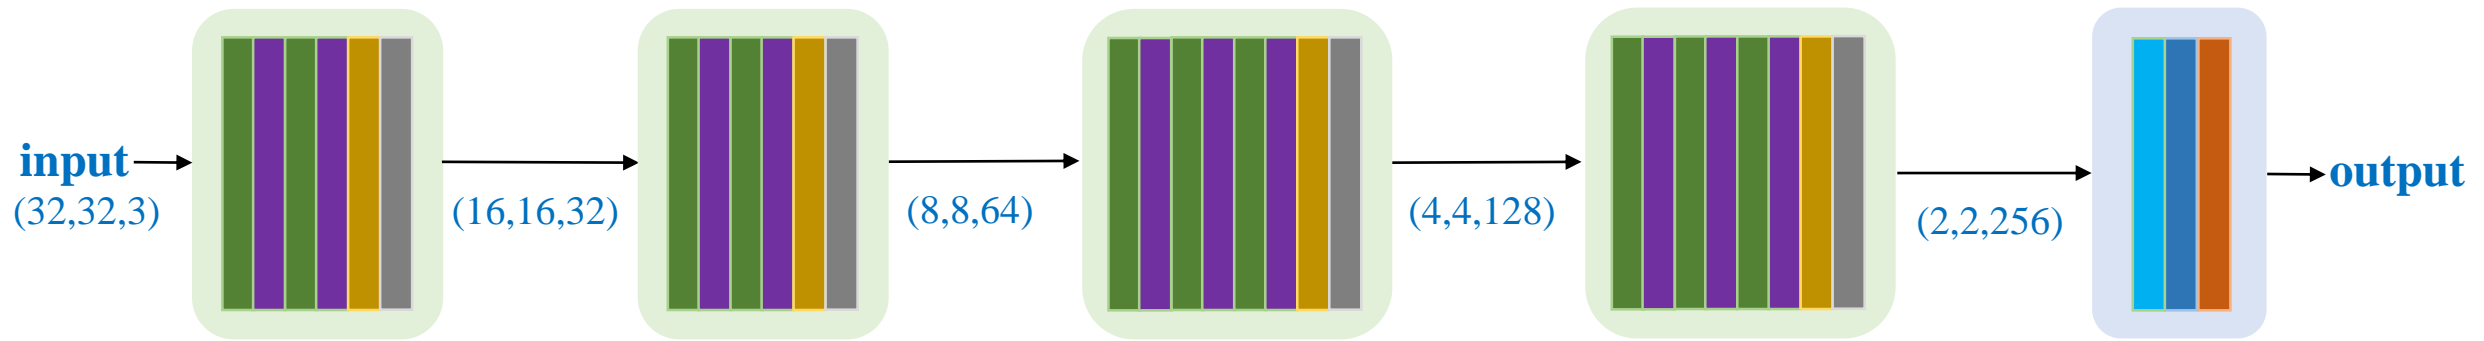
<https://arxiv.org/pdf/1803.08494.pdf>

“applying IN which does not only reduce the difference caused by domain changes, but also the illumination variation in single spectral images”

AFD-Net Aggregated Feature Difference Learning for Cross-Spectral Image Patch Matching (ICCV, 2019)

Model Generalization

❖ Trick 6: Instance normalization

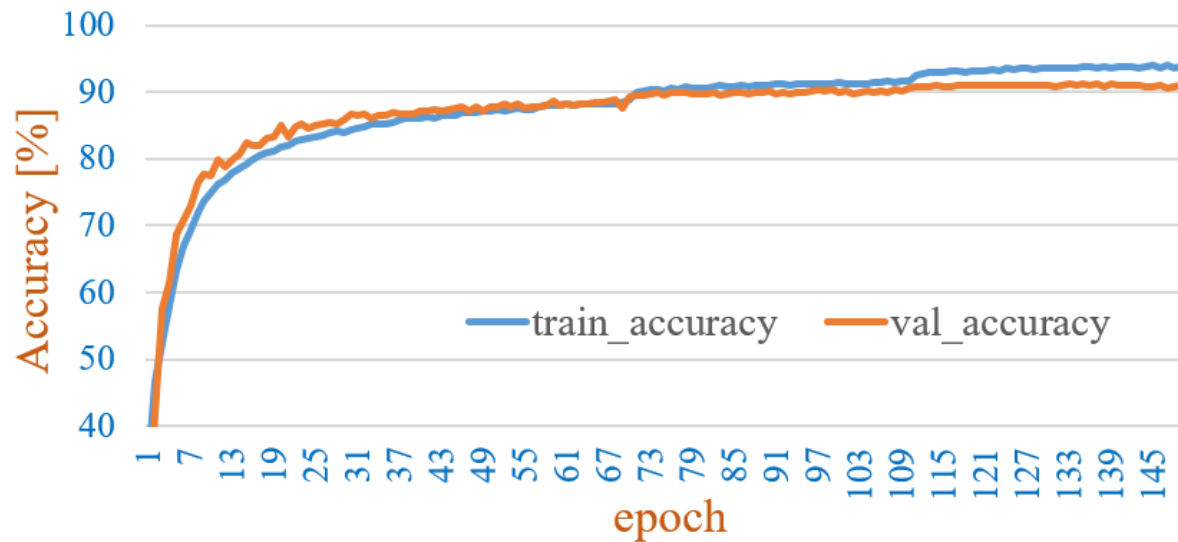


val_accuracy reaches to ~91.6%

Model Generalization

❖ Summary

Horizontal flip + crop-and-resize



val_accuracy reaches to ~91.6%

train_accuracy reaches to ~93.7%

Batch normalization

Dropout

Kernel regularization

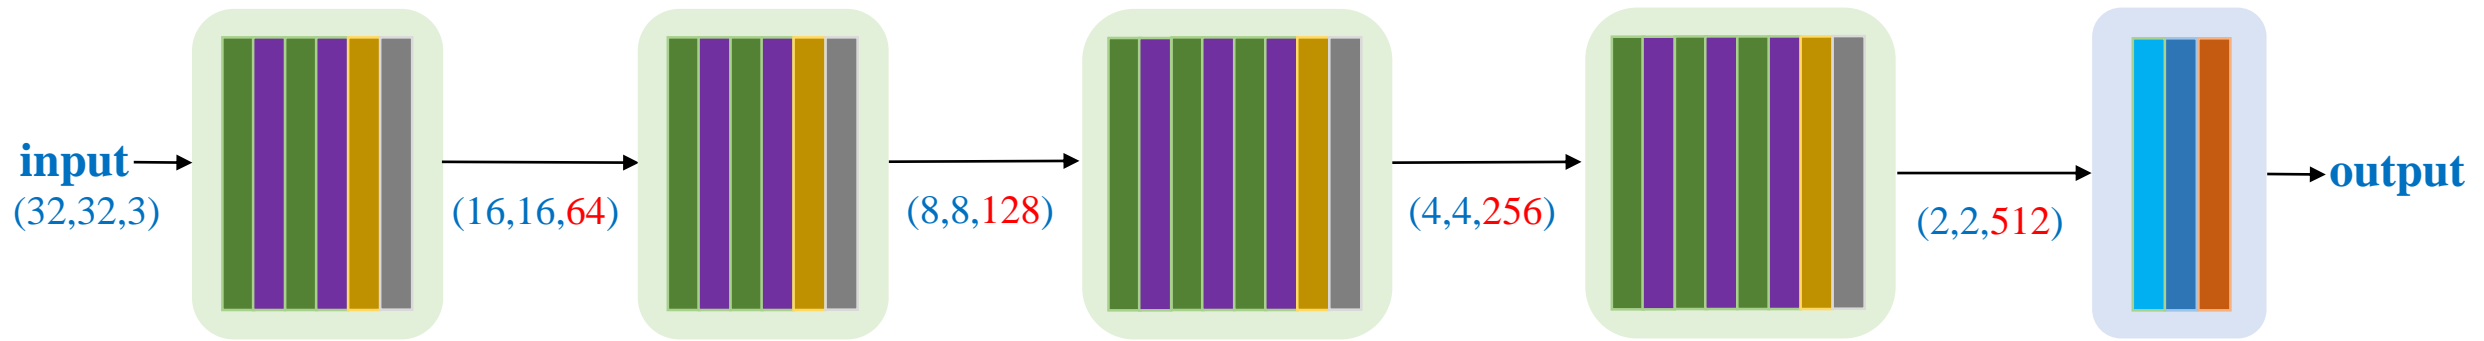
Data augmentation

Idea: try to increase train_accuracy,
expect val_accuracy increases too

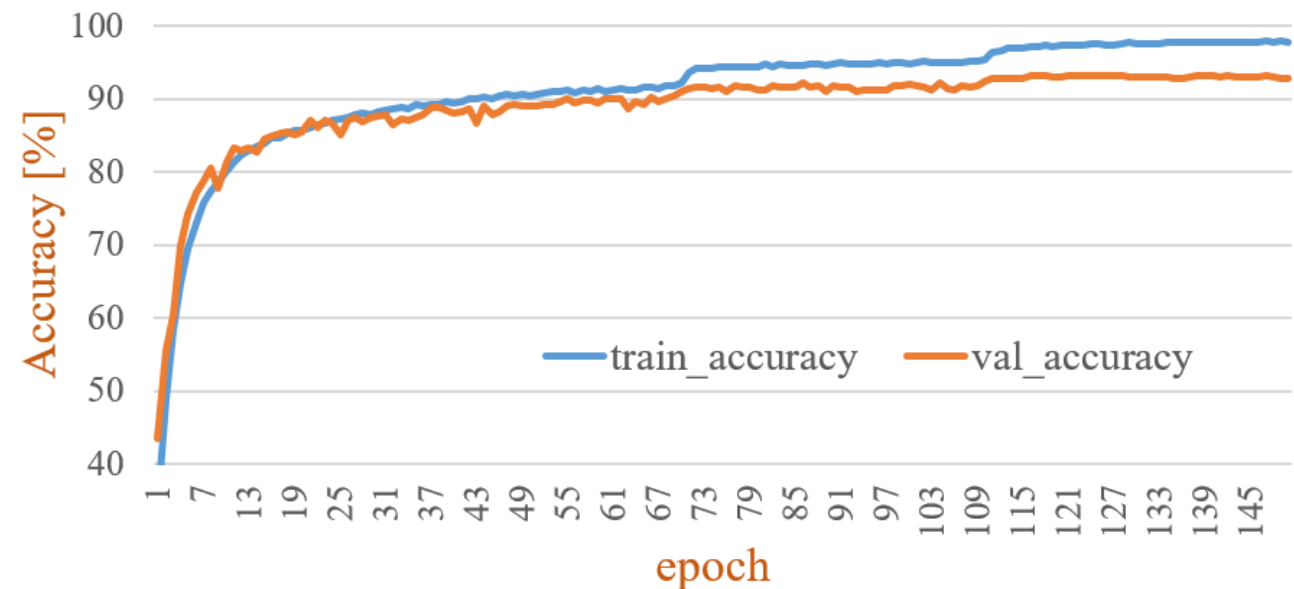
➡ Increase model capacity

Model Generalization

❖ Increase model capacity



val_accuracy reaches to ~93.6%
train_accuracy reaches to ~97.9%



Model Generalization



Summary

